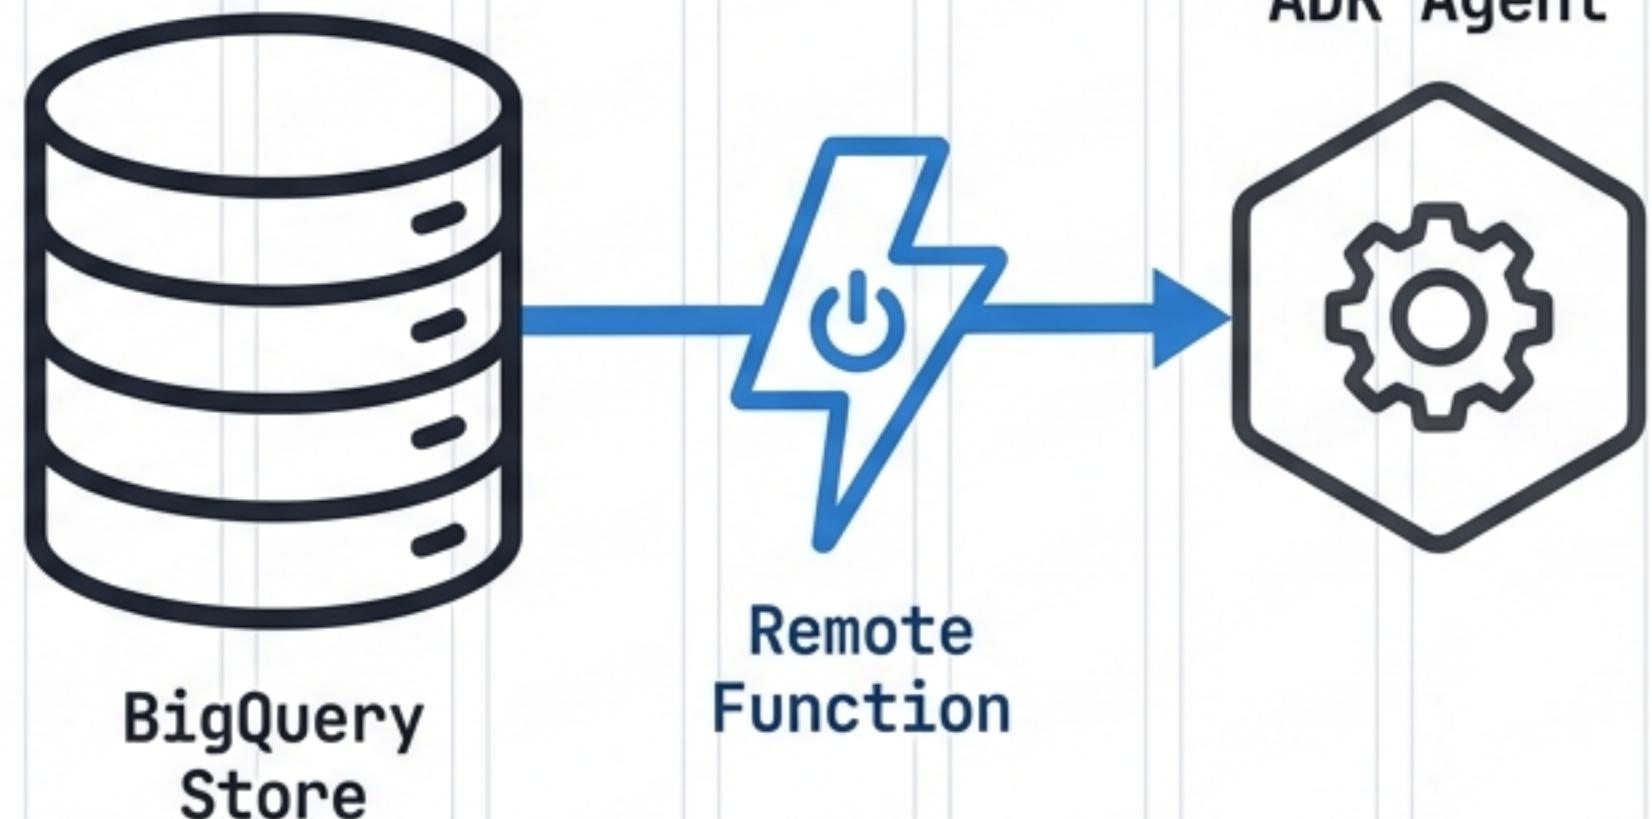


# Intelligent Scale: Processing 10k Records with BigQuery & ADK

An architectural review of connecting  
SQL-based data to Multi-Agent Systems  
using Remote Functions.

SYSTEM CONTEXT	
SCALE:	10,000+ Customer Records
STACK:	BigQuery / Cloud Run / Google ADK
PATTERN:	Remote Function Integration



# Executive Summary: The Strategy

## The Challenge



Run complex, multi-agent analysis (Security, Billing, Retention) on **10,000+** raw customer rows without **ETL** latency or data movement.

## The Solution



**BigQuery Remote Functions** act as a bridge to a **Cloud Run** service hosting Google's **Agent Development Kit (ADK)**. This brings compute directly to the data.

## The Verdict



Achieves '**Write Once, Run Anywhere**' logic with high development velocity. Main trade-offs involve session management overhead vs. raw **batch API** costs.

Result: Complex reasoning accessible via standard **SQL queries**.

# Moving Beyond Regex: The Need for Intelligent Analysis

## The Data Landscape

loyalty\_program (RECORD)

usage\_stats (RECORD)

security\_profile (RECORD)

billing\_summary (RECORD)

Problem: Standard SQL or Regex cannot reliably distinguish between nuanced states, such as a 'Billing Retry' vs. a 'Churn Risk'.

## The Bottleneck



Warehouse



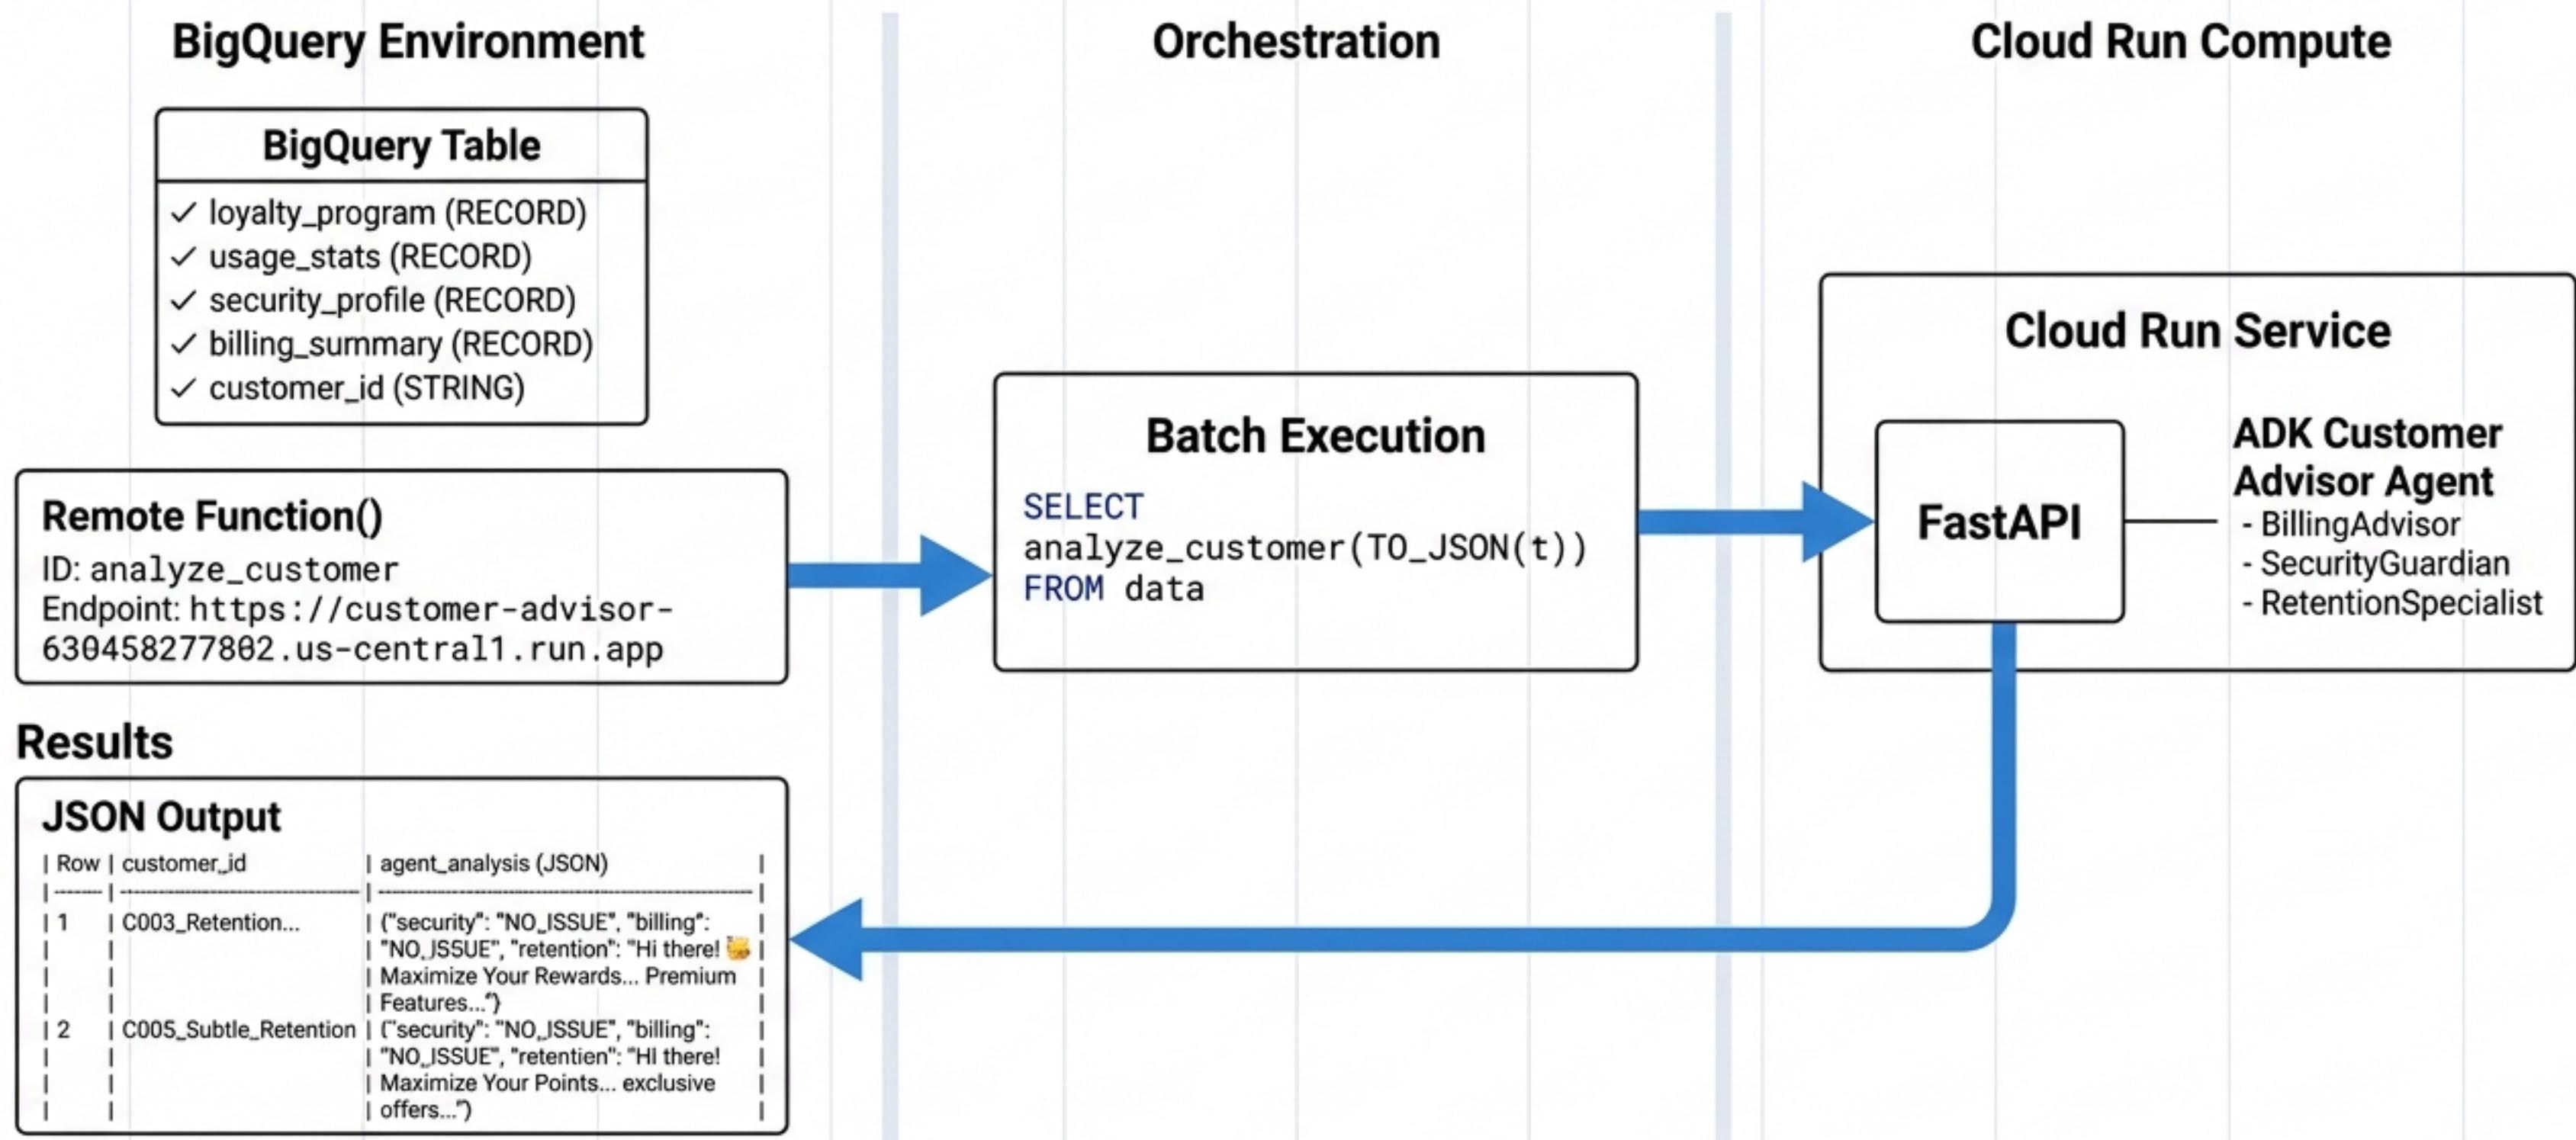
ETL / Data Movement



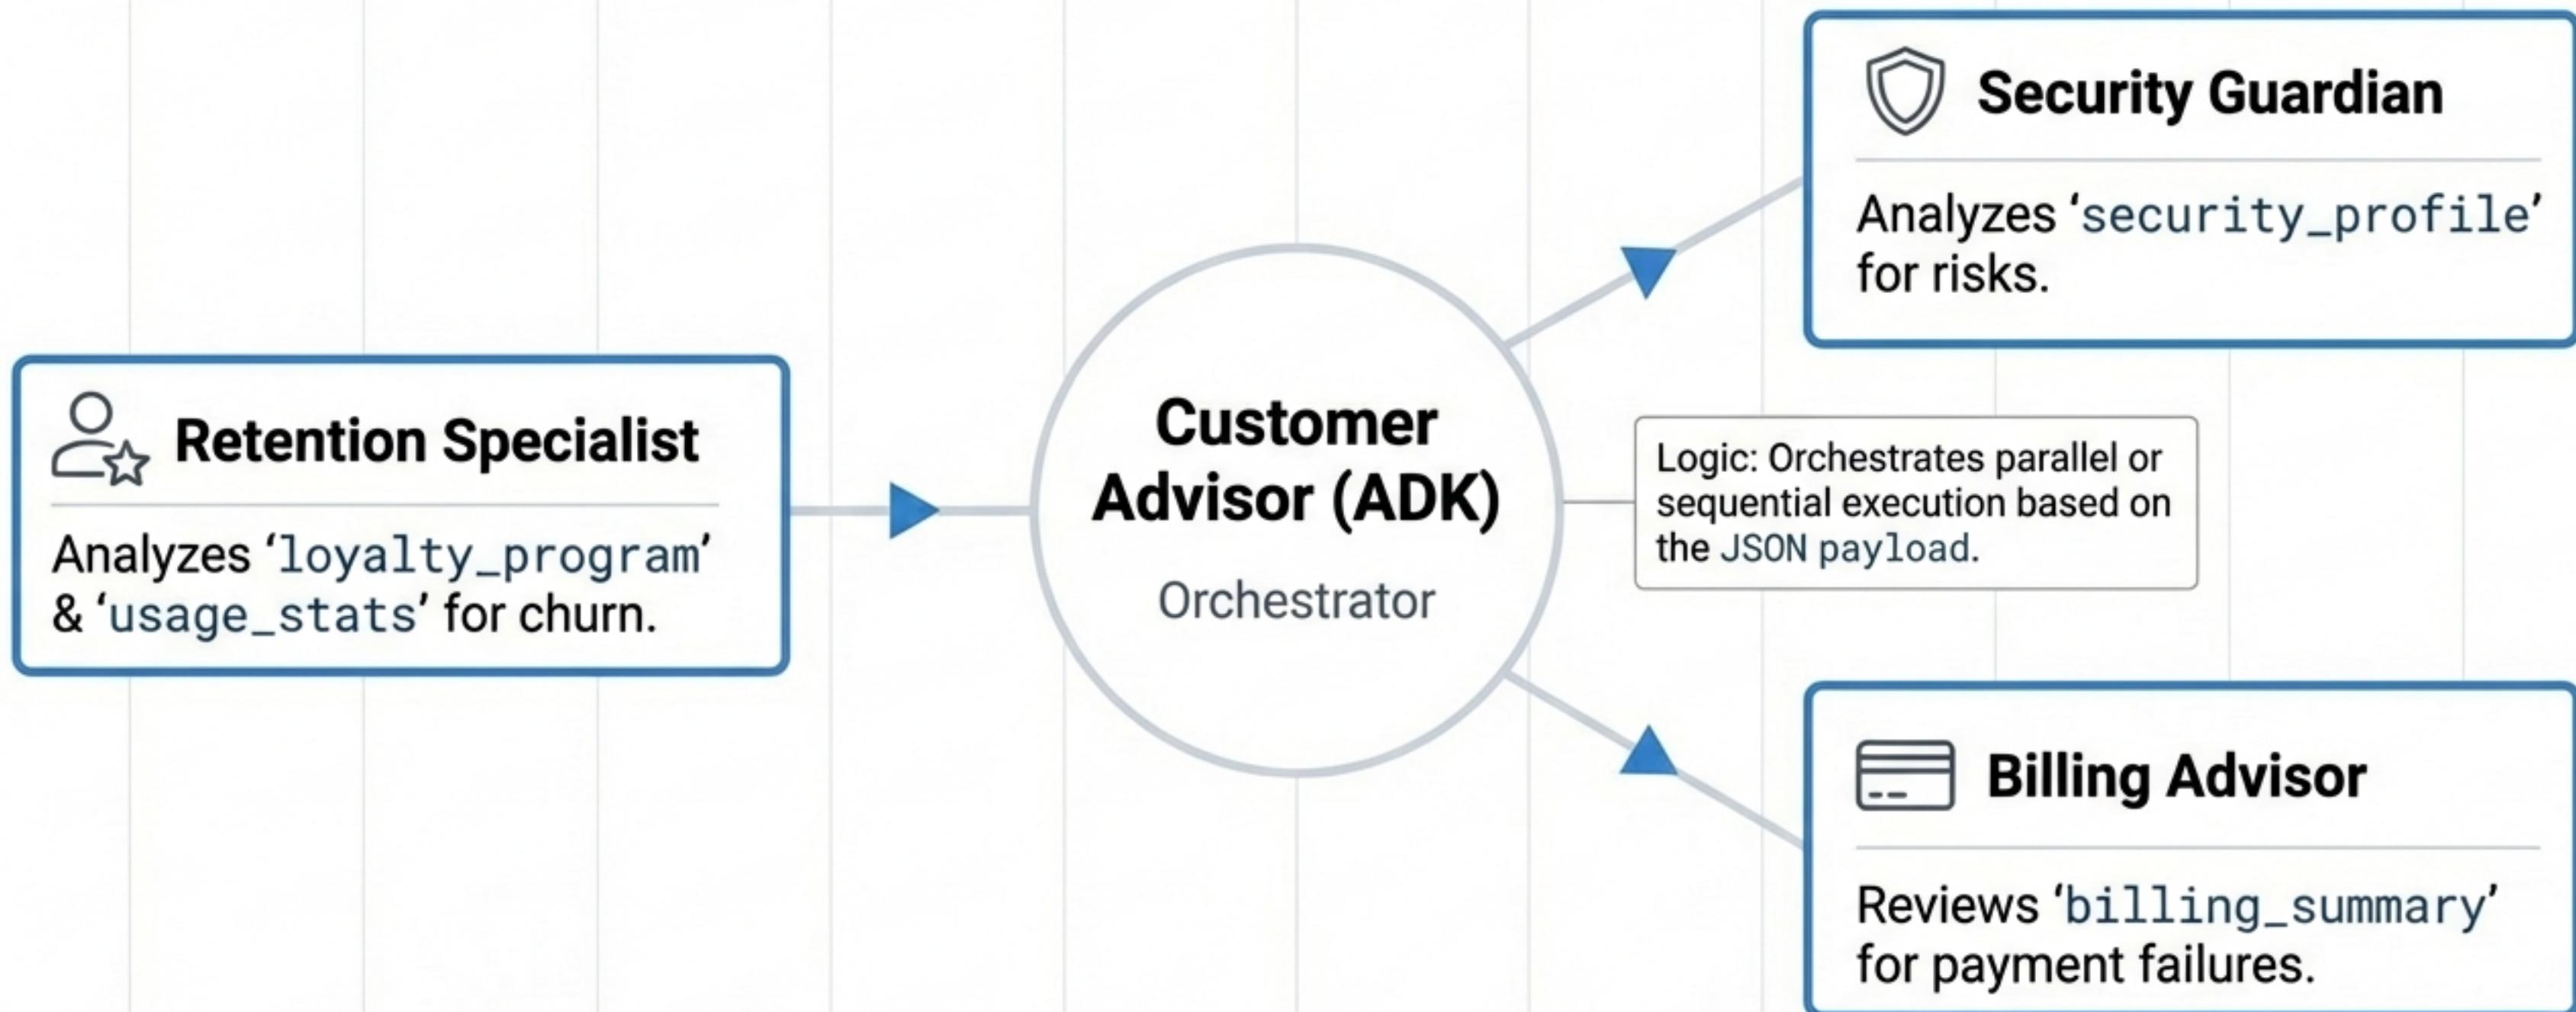
Processing Server

Constraint: Moving 10k records out of the Warehouse for processing adds latency and security risks. We need to bring the compute to the data.

# The Integration Pattern: Remote Functions



# The Brain: Multi-Agent Orchestration



# Invoking Intelligence via SQL

## SQL Query

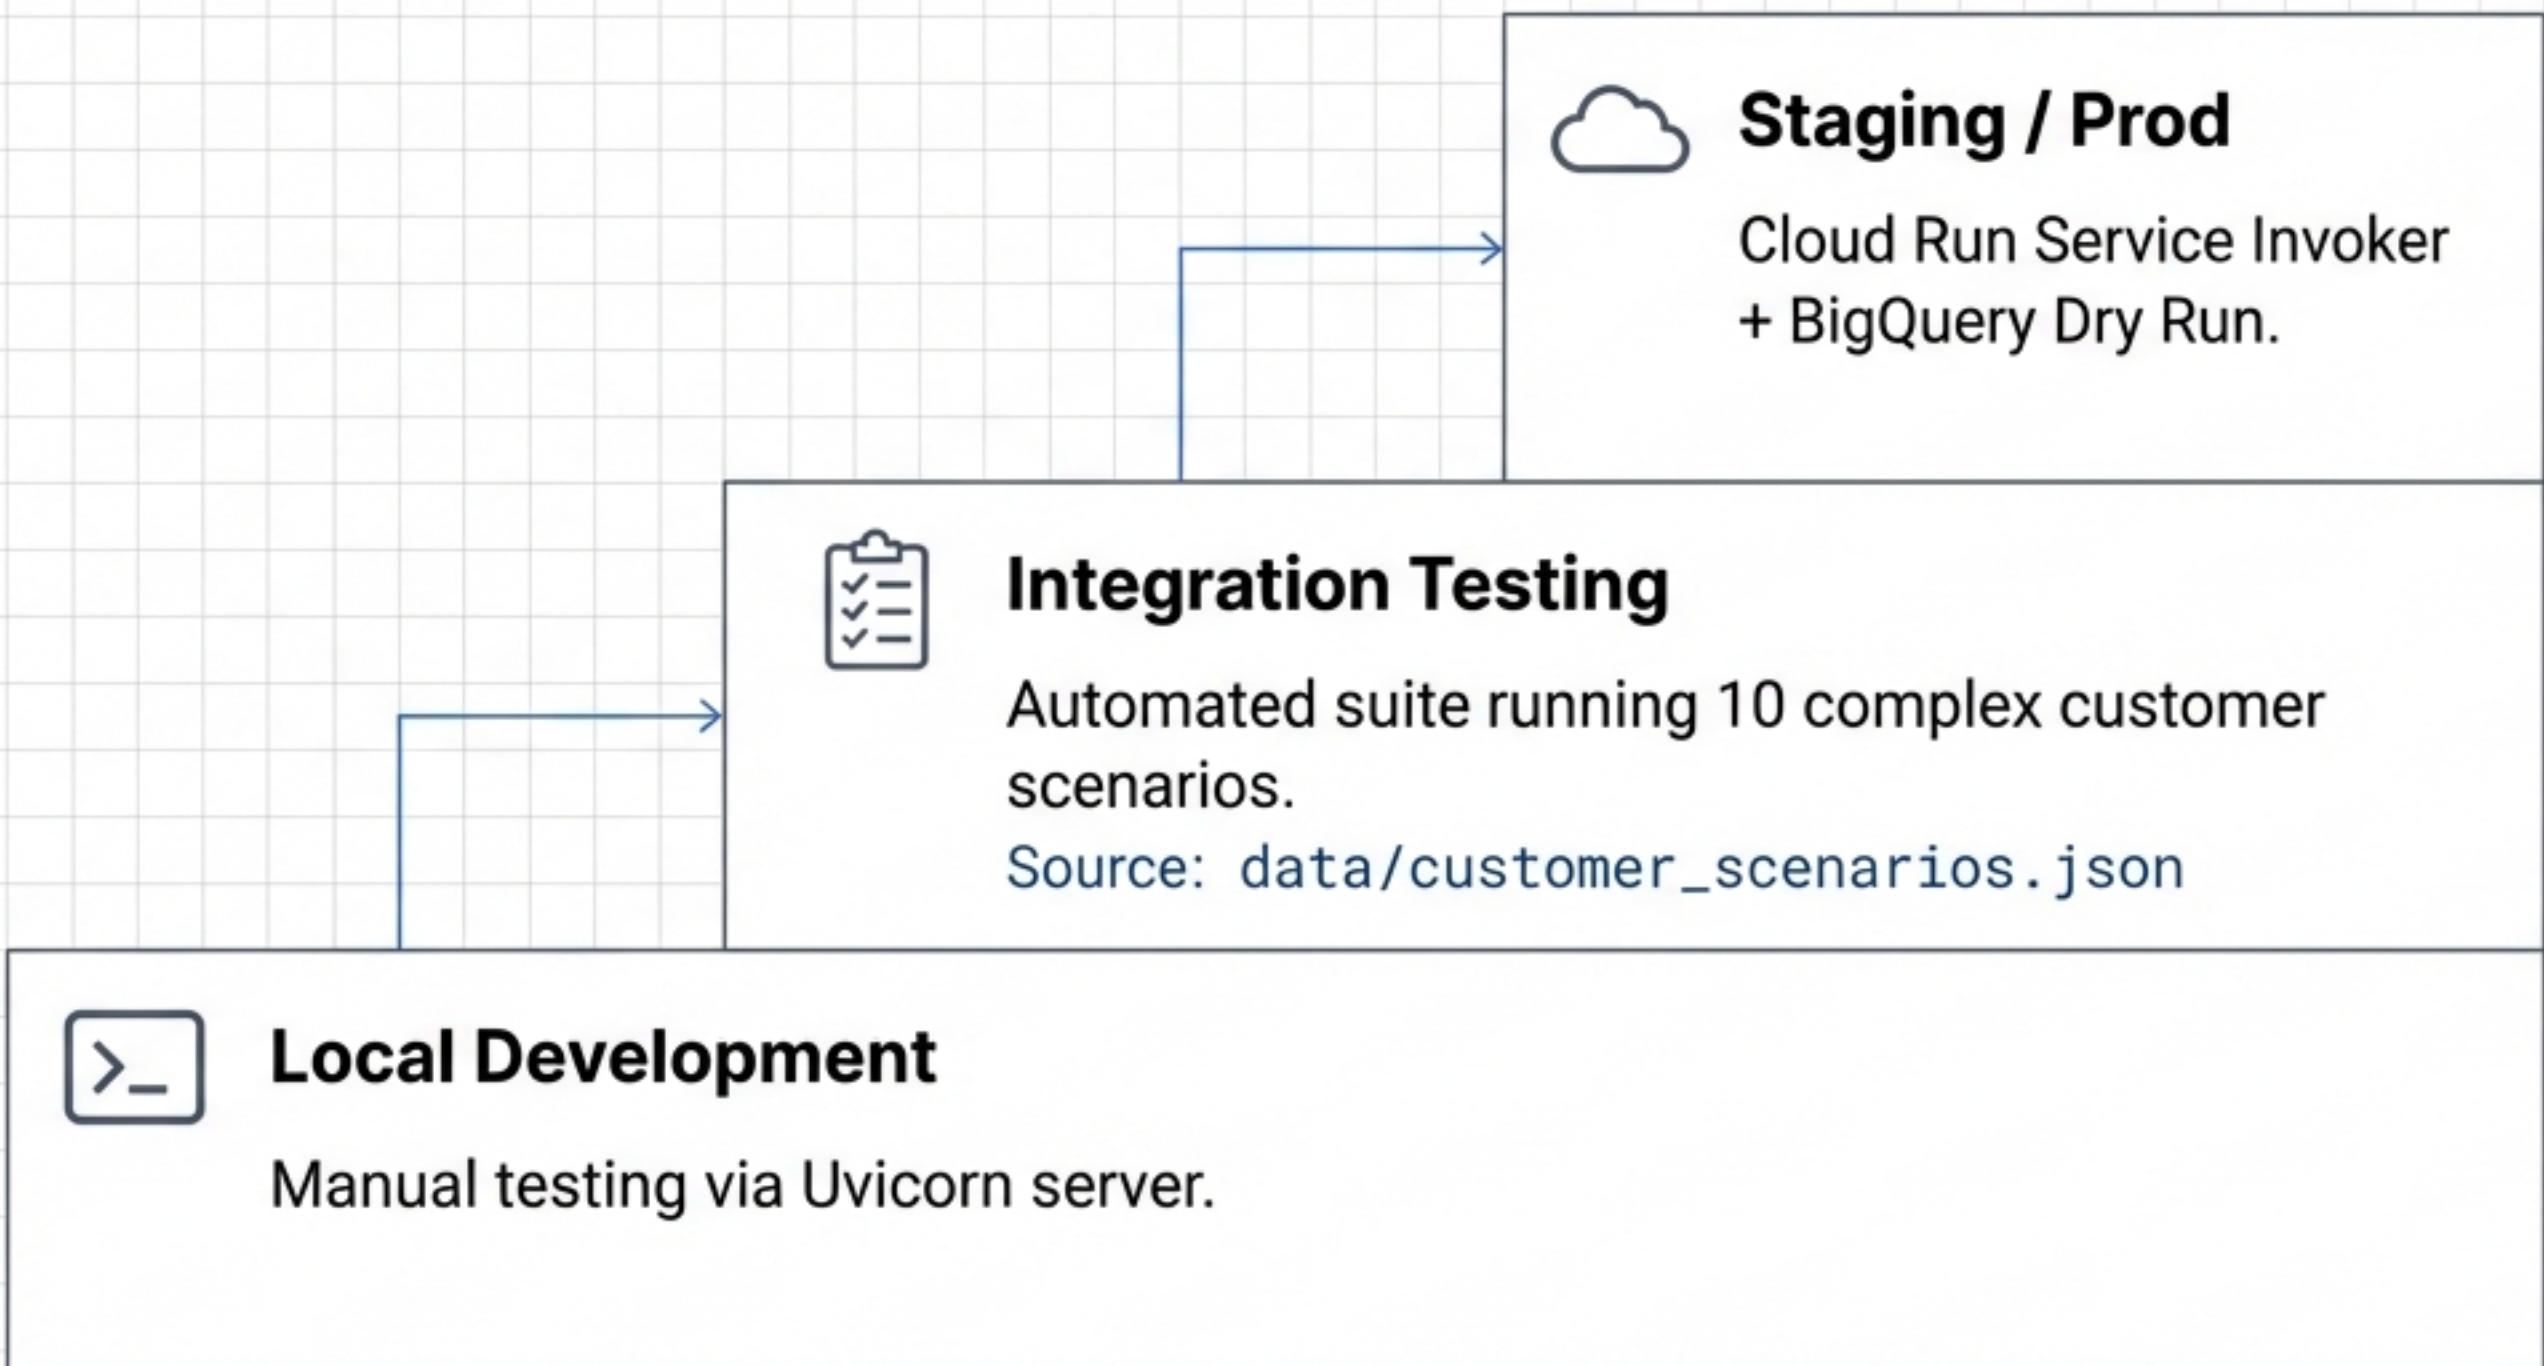
```
SELECT  
    customer_id,  
    remote_function_dataset.  
        analyze_customer(TO_JSON(t)) as  
            agent_analysis  
FROM  
    dataset.customer_scenarios AS t;
```

Encapsulated Complexity

## Structured Response

```
{  
    'security': 'NO_ISSUE',  
    'billing': 'NO_ISSUE',  
    'retention': 'Hi there! Maximize  
    Your Rewards. We notice you have  
    unused points...'}
```

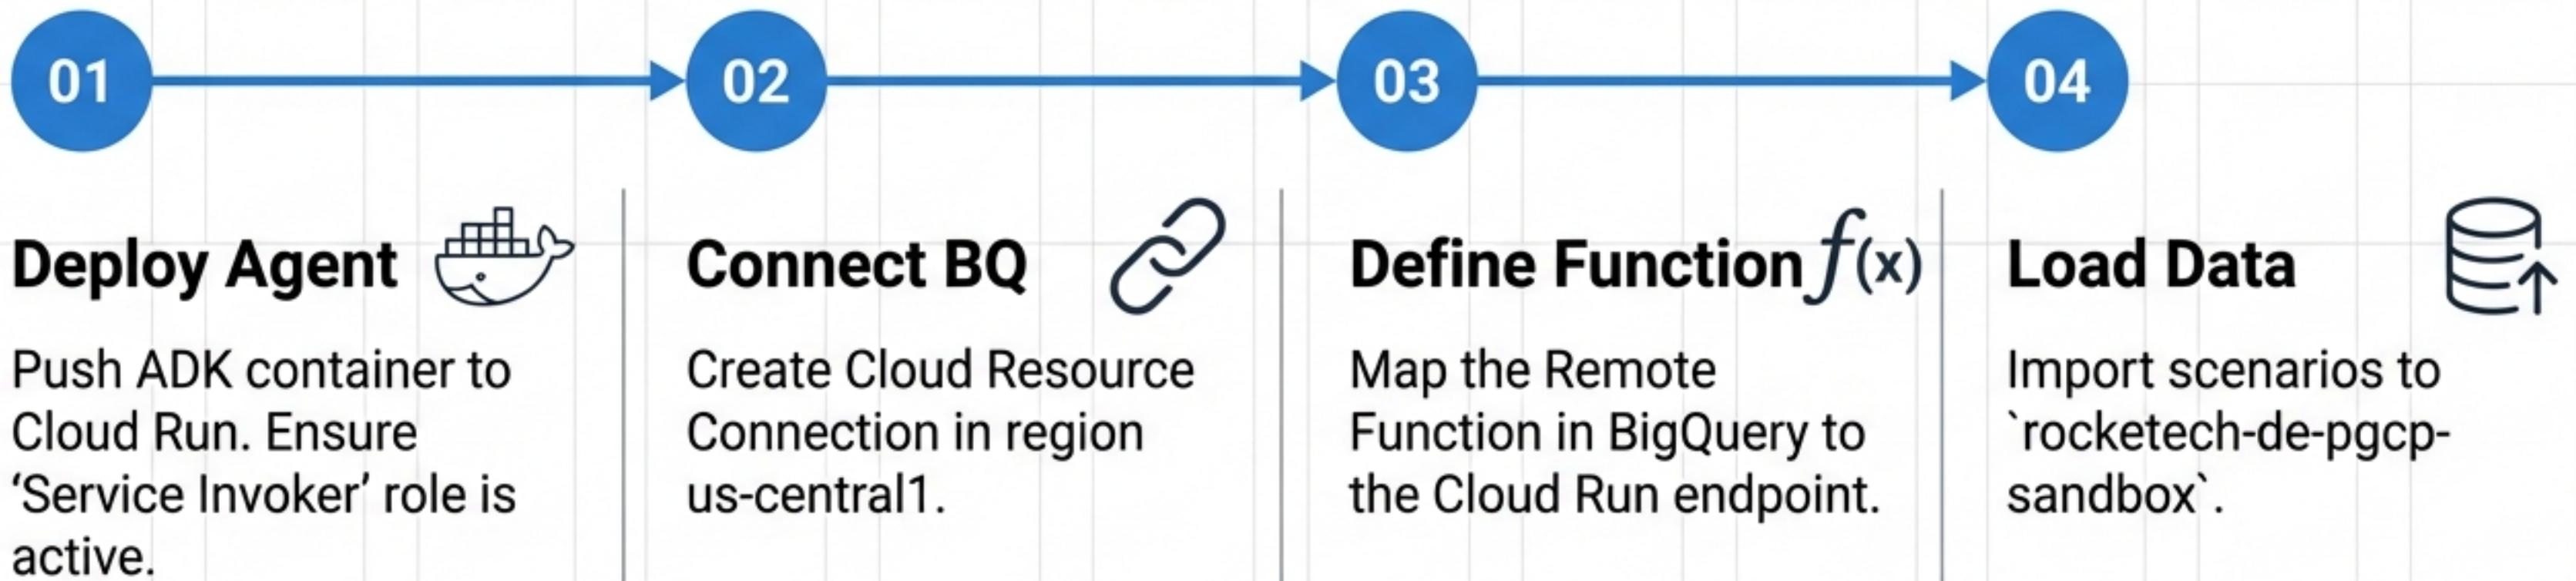
# Validating the Logic: Testing Strategy



## Validation Coverage:

- ✓ Security Logic
- ✓ Billing Logic
- Retention Logic
- Response Formatting

# Path to Production: Deployment & Configuration



**Accelerator:** Using the 'Agent Starter Pack' dramatically speeds up the prototype-to-production lifecycle.

# Architectural Verdict: ADK vs. Native Implementation

## The Wins (Why ADK)

 **Decoupling:** "Write Once, Use Anywhere". Same agent works in Gemini Enterprise or Event-Driven flows.

 **Sophistication:** Built-in Tracing, Tool calling, and Multi-Agent orchestration.

 **Velocity:** Faster dev cycles via the Agent Starter Pack.



## The Challenges

 **State Management:** ADK is session-based; adapting to stateless BQ functions requires memory optimization ([InMemorySessionService](#)).

 **Batching:** Modular sub-agents make standard batch inference harder/more costly than raw Vertex AI Batch API.

# Final Recommendation



**Primary Use Case:** Adopt this architecture for logic-heavy reasoning over warehouse data where development velocity and agent reusability are prioritized over raw batch throughput costs.

Alternative: Evaluate Vertex AI Batch API if volume > 1M records and logic is simple.

## Action Items

- 1. Development:** Use **Gemini CLI** for AI-assisted development (context pre-configured).

- 2. Testing:** Run the integration test suite to validate the 10 scenarios.